

Low-Level Design (LLD) for BigMart Sales Prediction Project

1. Introduction

This document provides a detailed technical design for building a machine learning model to predict sales for BigMart products across different outlets. It includes step-by-step information on the implementation, data handling, feature engineering, model training, and evaluation.

2. Input Specifications

2.1 Data Sources: - CSV file containing sales data for 2013 with product and store features.

2.2 Data Attributes: - **Product Features:** - Item_Identifier: Unique identifier for each product. - Item_Weight: Weight of the product. - Item_Fat_Content: Categorical feature indicating low or regular fat. - Item_Visibility: Percentage of total display area allocated to the product. - Item_Type: Category to which the product belongs. - Item_MRP: Maximum Retail Price of the product.

- **Store Features:**

- Outlet_Identifier: Unique identifier for each outlet.
- Outlet_Establishment_Year: The year when the outlet was established.
- Outlet_Size: Size of the outlet (small, medium, large).
- Outlet_Location_Type: Categorical feature indicating outlet location (urban, rural).
- Outlet_Type: Type of outlet (e.g., supermarket, grocery store).

- **Target Variable:**

- Item_Outlet_Sales: Sales for each product at a particular outlet.
-

3. Functional Requirements

3.1 Data Preprocessing: - **Objective:** Clean and preprocess data to make it suitable for machine learning.

Steps: - **Handling Missing Values:** - Item_Weight: Fill missing values using the median. - Outlet_Size: Fill missing values using the mode.

- **Outlier Detection:**

- Identify outliers in numerical features like Item_Visibility. Cap outliers above 99th percentile.

- **Feature Transformation:**

- Create Outlet_Age by subtracting Outlet_Establishment_Year from the current year (2013).
- Binning Item_Visibility into low, medium, and high categories to reduce skewness.

3.2 Exploratory Data Analysis (EDA): - Objective: Understand the distribution of data and relationships between variables.

Steps: - Univariate Analysis: - Plot histograms for Item_Outlet_Sales, Item_MRP, and Item_Weight. - Analyze categorical features using bar plots (Item_Type, Outlet_Type).

- **Bivariate Analysis:**

- Scatter plots to visualize Item_MRP vs. Item_Outlet_Sales.
- Box plots to check sales distribution across Outlet_Type and Outlet_Location_Type.
- Correlation heatmaps to find relationships between numerical variables.

Tools:

Use matplotlib and seaborn for visualizations.

3.3 Feature Engineering: - Objective: Create meaningful features that improve the model's performance.

Steps: - Create Derived Features: - Outlet_Age: Current year (2013) minus Outlet_Establishment_Year. - Item_Category: Group Item_Type into broader categories (e.g., Food, Non-Food).

- **Interaction Features:**

- Create interaction terms between Item_Type and Outlet_Type to capture product-outlet relationships.

Tools:

pandas and numpy.

3.4 Categorical Encoding: - Objective: Convert categorical variables into numeric form for modeling.

Steps: - Label Encoding: Use LabelEncoder for ordinal features such as Outlet_Size. - **One-Hot Encoding:** Apply OneHotEncoder or pd.get_dummies() on nominal variables like Item_Fat_Content, Outlet_Type, Item_Type, etc.

Tools:

sklearn.preprocessing.LabelEncoder and pandas.get_dummies().

3.5 Model Training: - Objective: Train various machine learning models to predict sales.

Steps: - Split Data: - Split data into training (70%) and testing (30%) sets using `train_test_split`.

- **Baseline Model:**

- **Linear Regression:**

- Train a basic linear regression model using `sklearn.linear_model.LinearRegression`.
 - Evaluate using `mean_absolute_error` and `r2_score`.

- **Regularized Models:**

- **Ridge and Lasso Regression:** Train Ridge and Lasso regression models with hyperparameter tuning.

- **Non-linear Models:**

- **Random Forest:** Train a `RandomForestRegressor` to capture non-linear relationships and feature interactions.
 - **XGBoost:** Train an `XGBRegressor` for better performance with optimized hyperparameters.

Tools:

`sklearn.linear_model`, `sklearn.ensemble.RandomForestRegressor`, `xgboost.XGBRegressor`.

3.6 Model Evaluation: - Objective: Evaluate the model's performance on test data.

Steps: - Calculate metrics: - **Mean Absolute Error (MAE):** Use `mean_absolute_error` to measure prediction accuracy. - **R-squared (R²):** Use `r2_score` to explain the variance captured by the model.

- Perform cross-validation to assess model stability.

Tools:

`sklearn.metrics.mean_absolute_error`, `sklearn.metrics.r2_score`, `sklearn.model_selection.cross_val_score`.

3.7 Model Tuning: - Objective: Optimize model hyperparameters to improve performance.

Steps: - **GridSearchCV:** Use grid search to tune hyperparameters such as `n_estimators` and `max_depth` for `RandomForest` and `XGBoost`. - **Feature Selection:** Apply feature importance analysis (e.g., based on `Random Forest` or `XGBoost` results) to eliminate less relevant features.

Tools:

`sklearn.model_selection.GridSearchCV`.

3.8 Final Model Selection: - Objective: Select the best-performing model based on evaluation metrics.

Steps: - Compare all models (Linear Regression, Ridge, Lasso, RandomForest, XGBoost) based on MAE and R². - Select the model with the lowest error and highest R² score for final deployment.

4. Non-Functional Requirements

4.1 Performance: - The system should be capable of predicting sales within a reasonable time (less than 2 seconds for each prediction). - The model training should complete within a few minutes for the given dataset.

4.2 Scalability: - The solution should handle increased data volume (e.g., more stores and products) without significant degradation in performance.

4.3 Usability: - The final model should be easy to deploy and use in future sales forecasting tasks.

5. Deployment Details

5.1 Model Saving: - Save the final model using joblib or pickle for future use.

Example: `python import joblib joblib.dump(best_model, 'final_sales_model.pkl')`

5.2 Model Inference: - Load the saved model and use it to predict sales for new data.

Example: `python model = joblib.load('final_sales_model.pkl') predictions = model.predict(new_data)`

6. Risks and Mitigation

Risk	Mitigation
Missing or inconsistent data	Use imputation and data validation techniques
Model overfitting	Use regularization (Ridge, Lasso) and cross-validation

Risk	Mitigation
Long training times for large datasets	Use optimized algorithms like XGBoost and RandomForest
Poor model performance	Perform extensive feature engineering and hyperparameter tuning

7. Testing and Validation

7.1 Unit Testing: - Test individual components like data preprocessing, feature engineering, and model training using unit tests.

7.2 Integration Testing: - Ensure that the data preprocessing pipeline and model training work cohesively by testing the full workflow.

7.3 Performance Testing: - Evaluate the model's performance using validation data and ensure that the prediction time is within acceptable limits.

8. Conclusion

The BigMart Sales Prediction project involves building a robust machine learning model to predict product sales across stores. This Low-Level Design document outlines each step involved in data preprocessing, feature engineering, model training, evaluation, and deployment. By following this plan, the project will deliver a scalable, efficient solution for forecasting sales.